

Week 1 — Facilitator Run-Sheet

Foundations & First Tool. A 10:00–16:00 day for five senior Sandvik GCC leaders. Built on one rule: teach a little, do a lot.

THE ROOM

MD · Finance Head · HR Head · HR Deputy · (one more functional lead). First structured course most have taken in 10+ years. Posture for the day: **calm expert who respects their seniority** — not high-energy L&D facilitation. Energy comes from the content landing, not from working the room.

DAY 1 PROMISE (REALISTIC) By 16:00 each leader will have **found the real problem they'll solve** and **architected & specced the agent that solves it** — sharp enough to hand to a builder. The build happens across Weeks 2-3. Day 1 is not "walk out with a working agent."

The first 15 minutes

Respect → a quiet measurement → a demonstration that proves the promise → the plan. No icebreakers, no naming their discomfort.

10:00 0-3 min	Bridge from March — evaluator to builder Plainly, no theatrics. Connect to the DRIVES session they already did. "In March we spent a session on how to <i>judge</i> AI — DRIVES, how to look at a proposal and decide if it's worth backing. That was the right place to start. But there's a ceiling to leading something you can't do yourself. Today we start removing that ceiling. Over these four weeks you'll stop being only the person who approves AI work — and become the person who can build it, and direct others building it."
10:03 3-6 min	The confidence read — a baseline, not an icebreaker Ask for one honest number. No discussion, no word cloud. The restraint is the point — you're measuring, like an engineer. "One question before we start, and I want an honest number. On a scale of 1 to 5 — how confident are you, right now, that you could take a messy problem from your own work and turn it into a precise, buildable spec for an AI agent? Hold that number. We come back to it at 4 o'clock." Why this question: it matches what the day actually delivers (a spec, not a built agent), so the number will genuinely move by close.
10:06 6-11 min	The black box — shown, not told LIVE DEMO The heart of the open. Use the finance vendor-quote demo (full data below). The 2026 point, not the 2024 one: with a strong model the vague prompt already gives a <i>good</i> answer — so the lesson is not "add detail." It's that the real skill is directing the reasoning: giving the model a method, a posture, and making it challenge itself. 1 Paste the vague prompt. It will give a competent answer — let that land. Say: "That's a good answer. The model is smart. So why are we here for four weeks?" 2 Paste the directed prompt — same three quotes, but it imposes a method (build the cost model), a posture (sceptical CFO hunting what's hidden), a self-challenge (argue the opposite), and forces a single committed call. 3 Let the difference land in silence: a worked normalisation model, the hidden clauses named, a devil's-advocate case, one decision. None of which the vague answer produced. "Both prompts went to the same model. The first asked it for an answer. The second told it how to think — gave it a method, made it hunt for what's hidden, made it argue against itself. That's the shift: two years ago the skill was being specific. Now the model is strong enough that the skill is directing its reasoning. That's the difference between using AI and being skilled with it — and it's the next four weeks."
10:11 11-15 min	The map + how today runs Now the three acts. Flag the hands-on beats. One line on the pre-work. "Three acts. Act 1, prompting — the everyday control you'll use most. Act 2, the engine — what's actually happening under the hood, so it stops being a black box. Act 3, RAG, agents, and the problem you'll carry for four weeks. The pattern all day is the same: I show you something, then you do it on a real problem of your own. You'll be building within the hour. You all played with these ideas in the pre-work this week — today you put them to work."

The full day · 10:00-16:00

Three real breaks + a full lunch. Hands-on blocks (tagged) are deliberately longer than teach blocks — that ratio is the whole point for this audience.

10:00 15 min	Open Bridge → confidence read → black-box demo → the map (detailed above).
10:15 60 min	Act 1 · Prompting HANDS-ON Teach (20): Anatomy of a production prompt · the controlled experiment (add one element per round, read what only that changed). Do (30): Build & bank 3 prompts on their own work. Trainer + floater circulate. Share (10): 2-3 read out; name the patterns & failure modes.
11:15 15 min	☕ Tea
11:30 55 min	Act 2 · The Engine Teach (25): How it predicts the next word · temperature · context window — just enough to demystify, not a CS lecture. Teach + call-and-response (30): Two things a leader decides (model choice + the dials) · the stack: compose, don't climb · "the cheapest layer that works wins." Quick round-the-room: "what does <i>this</i> task actually need?"
12:25 60 min	🍴 Lunch
13:25 50 min	Act 3a · RAG & Agents HANDS-ON Teach (15): RAG — flip retrieval on/off, the data rule (what can & can't go in). Teach (10): Agent vs chatbot — the think → act → observe loop. Do (25): "AI interviews you" — each leader runs the capstone-finder prompt in Claude and surfaces their real problem.
14:15 45 min	Act 3b · Architect your problem HANDS-ON Do (25): Make the machine architect their problem across the stack; then push every layer down to the cheapest that works. Do (20): Out-judge the machine — prove its first map wrong, stress-test the riskiest assumption.
15:00 15 min	☕ Tea
15:15 30 min	Act 3c · Write & pressure-test the spec HANDS-ON Do (30): The three-part test, then write the agent spec. Pressure-test it with Claude, then with a peer in the room.

15:45
15 min

Close — prove the shift

Re-ask the 1–5 confidence number; show the move live. Each leader says their capstone problem in **one sentence** — public commitment in front of peers. Set up the Week 1 homework ("Map the Stack") and what Week 2 does with it.

Coverage rule: during every hands-on block, a second person floats. When a senior leader gets stuck and no one reaches them within ~60 seconds, they quietly disengage and decide "this isn't for me." Coverage is non-negotiable.

Demo data · Finance vendor-quote reconciliation

Context: Sandvik's Pune technology GCC builds & supports engineering software (Cimatron, Mastercam, GibbsCAM, SigmaNest, Metrologic) used by Sandvik customers globally. The Finance Head is evaluating three vendors for an **annual managed cloud-testing & CI/CD infrastructure contract** for the dev teams. The three quotes are deliberately structured differently — different units, different inclusions, hidden assumptions. This is exactly the mess a finance head untangles by hand.

Before Friday: rehearse both prompts live on Opus (or whichever model you'll use) and confirm the gap is actually visible on screen. With a strong model the vague prompt gives a decent answer — if the two land too close, make the vague one lazier ("which is cheapest?") so it anchors on headline numbers. Screenshot the directed output as a wifi-failure backup. Attach or paste the three vendor PDFs for both prompts.

PROMPT 1 — THE VAGUE VERSION (PASTE FIRST)

Here are three vendor proposals for our build/test infrastructure. Which is the best deal?

[paste the three vendor PDFs, or the three summary lines below]
Vendor A: \$18,000/month, managed, includes support.
Vendor B: \$190,000/year all-in, 20% off if signed by quarter-end.
Vendor C: \$9.50 per build-hour, pay as you go, plus \$4,000/month.

On Opus this gives a *competent* answer — it'll do arithmetic and probably note the quotes aren't directly comparable. That's the point: don't pretend it failed. Say "that's a good answer — so why are we here for four weeks?" Then show what directing the reasoning adds.

PROMPT 2 — THE DIRECTED VERSION (PASTE SECOND)

Same model, same quotes. What changes is that it imposes **a method, a posture, a self-challenge, and a forced commitment** — four things Opus will not do unprompted.

You are a sceptical CFO-level advisor to the Finance Head of a 500-engineer captive technology GCC. Assume each of these three proposals has been deliberately structured to look good on paper and to hide its real cost or risk. Your job is to find what each one is hiding. Don't trust any headline number.

The three proposals are attached. Our real situation:
- ~6,200 build-hours/month, spiking to ~8,500 in release weeks (4 release weeks/year).
- We require 24x7 support with a 1-hour P1 SLA.
- We need 3 regions (India, EU, US) for data residency.
- Current annual spend on what this replaces: ~\$210,000.

Work in this exact order:

METHOD — Build one normalised cost model. Convert all three to a true 12-month total at OUR actual volume including release spikes. Show the model and your assumptions, not just the totals. Do not use their headline framing — restate everything in the same units.

HUNT — For each vendor, name the single specific clause engineered to win the comparison on paper, and what it actually costs us once priced in (auto-renew traps, undefined "fair use", excluded regions, overage rates, consumption exposure).

DECIDE — Give me exactly ONE recommendation and the single biggest risk of being wrong. No hedging, no "it depends". If you need an assumption, state it and move on.

CHALLENGE YOURSELF — Now argue the opposite case. What would have to be true for your rejected options to actually be the right call? Has that changed your recommendation?

QUESTION — One sharp question to put to each vendor before signing that would expose what they're hiding.

What the good output should surface (facilitator cheat-sheet)

Run it live, but know what's coming so you can narrate it. The real numbers below assume ~6,200 hrs/mo base, +2,300 extra hrs across 4 release weeks ≈ **~74,400 + ~9,200 = ~83,600 build-hours/year**.

Vendor	True 12-month cost at our usage	The hidden assumption they hope you miss
A · CloudForge	\$18k×12 = \$216k base, but only 5,000 hrs/mo included. We use ~6,200–8,500. Overage ≈ (avg ~1,800 extra hrs/mo ×12 ≈ 21,600 hrs) × \$3.20 ≈ +\$69k . Plus 2 extra regions ×\$2,500×12 = +\$60k . ≈ \$345k	"Fully managed" hides a low included-hours cap, expensive overage, single region, and a 90-day auto-renew trap.
B · TestPipe	\$190k/yr, or \$152k if you accept a 24-month lock . But 3rd region (US) is excluded and quoted separately — likely +\$30–50k . So real year-1 ≈ \$190k–\$240k , and the "discount" trades cash for a 2-year commitment.	"Unlimited build-hours" undercut by an undefined "fair use policy." The 20% discount is bait for a 24-month lock. 4-hr P1 SLA misses our 1-hr need.
C · BuildScale	Base 74,400 hrs: months under 7,000 hrs at \$9.50, release months above at \$7.50. Blended ≈ \$9.30 × 83,600 ≈ \$778k in build-hours + \$48k platform . ≈ \$826k — unless usage is far lower than stated.	Only one that meets the 1-hr SLA and all 3 regions out of the box — but consumption pricing is wildly exposed to our actual hours . Cheapest on paper per unit, most expensive in total at our volume.

Wins on cost
Vendor B (~\$190–240k) — closest to current \$210k spend, only one near a flat predictable number once you add the missing region.

Wins on risk / fit
Vendor C on SLA + regions + no lock-in — but its cost is the most volatile. The disagreement: cheapest-to-run (B) carries a 24-month commitment; best-fit (C) carries open-ended cost exposure.

The one question for each vendor

Vendor	Ask before signing
A	"Cap our all-in monthly at a fixed number including overage and all 3 regions — what is it?" (Forces the true price into the open.)
B	"Define 'fair use' in writing as a build-hour number, give me the 3-region all-in price, and give me the 20% on a 12-month term." (Removes the lock + the undefined ceiling.)
C	"Give me a monthly cap / committed-use rate so we're not exposed on release months." (Converts open consumption into a bounded number.)

The teaching beat to say out loud after the reveal: "The vague prompt already gave a fair answer — the model is strong. But look at what the second one produced that the first didn't: a worked cost model, the hidden clause in each quote named, a case argued against itself, and one committed decision. I didn't give it more facts — I gave it a way to think. A method, a sceptical posture, and an instruction to challenge its own conclusion. That's the skill that's scarce now. Not asking AI better — directing how it reasons."

people equation · AI Capability Program · Week 1 facilitator run-sheet · confidential
Synthetic vendor data for training demonstration only — figures illustrative, tuned to a Sandvik-style technology GCC.